

CODEBUDDY: A Natural Language Code Explanation Generator Using Fine-Tuned CODET5 and Retrieval-Augmented Generation for Enhanced Developer Productivity

Vinay Jogani¹, Shruti Dhamdhere¹

¹College Of Engineering
Northeastern University
Boston, MA, USA

August 11, 2025

Abstract

The exponential growth in software complexity has created critical challenges in code comprehension, with developers spending approximately 60% of their time understanding existing codebases rather than implementing new features. Current automated documentation solutions suffer from limitations including lack of contextual awareness, generic explanations, and inability to engage in interactive dialogue.

This paper presents CODEBUDDY, a novel AI-powered system that generates natural language explanations for Python code through the integration of fine-tuned CODET5 with Retrieval-Augmented Generation (RAG). Our approach combines specialized transformer architecture for code understanding with vector database-driven retrieval to provide context-aware explanations and interactive question-answering capabilities.

CODEBUDDY was trained on a comprehensive multi-dataset corpus including CodeSearchNet (2M+ Python functions), DocString Dataset, HumanEval, and MBPP, totaling over 2.1 million code-documentation pairs. The system architecture integrates CODET5 encoding with FAISS vector storage and a RAG pipeline for contextual retrieval and generation.

Extensive evaluation demonstrates significant improvements over existing approaches: BLEU score of 0.72 (vs. 0.58 GPT-4 baseline, +24.1%), BERTScore of 0.84, and Exact Match accuracy of 78% for interactive Q&A. The system maintains practical response times under 500ms while providing accurate explanations across code complexity levels.

A production-ready Streamlit interface enables three core functionalities: instant code explanation, interactive questioning, and semantic code search. Ablation studies confirm the critical contributions of domain-specific fine-tuning (+18% BLEU improvement) and RAG integration (+50% Q&A accuracy improvement). This work addresses fundamental gaps in automated code documentation by providing specialized, context-aware explanations that significantly enhance developer productivity and code comprehension capabilities.

Keywords: Code explanation, Natural language generation, CodeT5, Retrieval-augmented generation, Software documentation, Code comprehension, Transformer models, Python analysis

1 Introduction

1.1 Motivation and Background

Software development has evolved into an increasingly complex discipline where understanding existing code often represents the primary bottleneck in developer productivity. Empirical studies in software engineering demonstrate that developers allocate approximately 58-70% of their working time to program comprehension activities rather than new code implementation [10, 19].

This comprehension burden becomes particularly acute during several critical development scenarios: onboarding new team members to existing codebases, conducting thorough code reviews, maintaining legacy systems with insufficient documentation, and integrating third-party libraries or frameworks.

The traditional paradigm of manual code documentation has proven inadequate for modern software development demands. Manual documentation suffers from several fundamental limitations: it requires significant time investment that competes directly with feature development priorities, documentation quality varies dramatically based on individual developer expertise and motivation, maintaining documentation currency with rapidly evolving codebases proves challenging, and comprehensive documentation of complex algorithmic logic often exceeds practical time constraints.

Contemporary software projects frequently involve codebases containing millions of lines of code across hundreds of files, making comprehensive manual documentation practically infeasible. The situation is further complicated by the increasing prevalence of machine learning systems, complex distributed architectures, and domain-specific libraries that require specialized knowledge for effective comprehension.

1.2 Problem Statement

The central research problem addressed by this work is the development of automated, intelligent systems capable of generating accurate, contextually-aware natural language explanations for source code. Current approaches to automated code explanation fall into several categories, each exhibiting significant limitations that prevent widespread adoption:

Static Analysis Tools: Traditional static analysis systems can extract structural information such as function signatures, variable declarations, and control flow patterns. However, these tools lack semantic understanding of code functionality and cannot generate meaningful natural language explanations that capture algorithmic logic or design intent.

Template-Based Documentation Generators: Systems like Doxygen and Javadoc can generate documentation from structured comments and code annotations. While useful for API documentation, these approaches require extensive manual annotation and cannot provide explanations for undocumented code or capture complex algorithmic behavior.

Generic Language Models: Recent advances in large language models (LLMs) such as GPT-4 have demonstrated impressive capabilities for code explanation tasks. However, these general-purpose models lack specialized training for code understanding and often produce explanations that miss domain-specific patterns, technical nuances, or contextual relationships within larger codebases.

Rule-Based Expert Systems: Some specialized tools attempt to explain code using predefined rules and pattern matching. These systems are limited by their rule sets and cannot adapt to novel code patterns or provide explanations for innovative algorithmic approaches.

1.3 Research Objectives and Hypotheses

This research aims to address the identified limitations through the development and evaluation of CODEBUDDY, an AI-powered code explanation system. The primary research objectives are:

Primary Objective: Design and implement an intelligent code explanation system that generates accurate, contextually-aware natural language descriptions of Python source code while supporting interactive question-answering capabilities.

Technical Objectives:

1. Develop a specialized transformer-based architecture optimized for code understanding tasks
2. Implement a retrieval-augmented generation system that leverages contextual code examples for enhanced explanation quality
3. Create a comprehensive evaluation framework incorporating multiple quantitative and qualitative assessment metrics

4. Build a production-ready system with user-friendly interface and deployment capabilities

Performance Hypotheses:

1. **H1:** Domain-specific fine-tuning of CODET5 will produce significantly higher quality code explanations compared to generic language models
2. **H2:** Integration of retrieval-augmented generation will improve contextual accuracy and enable effective interactive question-answering
3. **H3:** The combined system will maintain practical response times suitable for interactive development environments

1.4 Contributions

This work makes several significant contributions to the fields of automated software engineering and natural language processing:

Novel Technical Architecture: Integration of fine-tuned CODET5 with RAG pipeline represents a novel approach to code explanation. This combines specialized code understanding with contextual retrieval capabilities.

Comprehensive Evaluation Framework: Development of multi-metric evaluation approach incorporating BLEU scores, BERTScore semantic similarity, exact match accuracy, and response time analysis. This provides robust assessment methodology for code explanation systems.

Production-Ready Implementation: Complete system implementation including user interface, error handling, logging, and deployment capabilities. This demonstrates practical applicability beyond research prototypes.

Large-Scale Dataset Integration: Systematic combination and preprocessing of multiple established code documentation datasets. This creates comprehensive training corpus for specialized model development.

Empirical Performance Validation: Rigorous experimental evaluation demonstrating measurable improvements over existing baselines across multiple assessment dimensions.

2 Related Work

2.1 Automated Code Documentation

The field of automated code documentation has evolved significantly over the past two decades. Early approaches focused on extracting structural information from source code to generate basic documentation templates [17]. These systems could identify function signatures, parameter types, and return values but lacked the semantic understanding necessary for meaningful explanation generation.

Sridhara et al. [17] pioneered early work in automatic comment generation for Java methods, introducing the concept of extracting semantic information from code structure. Their approach used static analysis to identify code patterns and generate template-based comments, establishing important foundations for subsequent research.

More recent work has leveraged advances in natural language processing to improve documentation quality. Iyer et al. [7] introduced neural attention models for source code summarization, demonstrating that sequence-to-sequence architectures could capture more complex relationships between code structure and natural language descriptions.

2.2 Transformer Models for Code Understanding

The application of transformer architectures to code understanding tasks has shown remarkable progress. CodeBERT [4] was among the first models to demonstrate the effectiveness of BERT-style pretraining for code-related tasks, achieving significant improvements on code search and documentation generation benchmarks.

Wang et al. [18] introduced CODET5, which extends the T5 encoder-decoder architecture specifically for code understanding and generation tasks. CODET5 incorporates identifier-aware pretraining objectives that better capture code semantics compared to general-purpose language models. The model demonstrates superior performance across multiple code-related benchmarks including code summarization, code generation, and code translation tasks.

Subsequent work has explored various adaptations of transformer architectures for specialized code tasks. GraphCodeBERT [5] incorporates code structure information through graph neural networks, while CodeGPT [13] focuses on autoregressive generation for code completion and synthesis tasks.

2.3 Retrieval-Augmented Generation Systems

Retrieval-Augmented Generation emerged as a powerful paradigm for enhancing language model capabilities by incorporating external knowledge during generation [11]. The core insight is that retrieving relevant context from large knowledge bases can significantly improve generation quality, particularly for knowledge-intensive tasks.

In the context of code understanding, several researchers have explored RAG applications. Zheng et al. [21] investigated neural code explanation with retrieval augmentation, demonstrating improvements in explanation quality when relevant code examples are retrieved during generation. Their work established important foundations for combining retrieval systems with code-specific language models.

The integration of vector databases with transformer models has enabled efficient similarity search across large

code repositories. FAISS [8] provides scalable approximate nearest neighbor search capabilities that make real-time retrieval practical for interactive applications.

2.4 Evaluation Methodologies

Evaluating code explanation quality presents unique challenges that distinguish it from general natural language generation tasks. Traditional metrics like BLEU [15] and ROUGE [12] provide useful quantitative measures but may not capture technical accuracy or semantic correctness specific to code domains.

Zhang et al. [20] introduced BERTScore as a semantic similarity metric that better correlates with human judgment compared to n-gram based metrics. For code-specific tasks, CodeBLEU [16] adapts BLEU scoring to account for code structure and syntax.

Recent work has emphasized the importance of human evaluation for code explanation tasks [1]. Expert assessment of technical accuracy, completeness, and clarity provides essential validation that automated metrics may miss.

2.5 Research Gaps and Opportunities

Despite significant progress, existing approaches exhibit several limitations that create opportunities for improvement:

Limited Contextual Awareness: Most systems analyze code in isolation without considering broader codebase context, leading to explanations that miss important relationships and dependencies.

Generic Training Paradigms: Models trained on general text corpora lack deep understanding of programming idioms, algorithmic patterns, and domain-specific conventions.

Static Explanation Generation: Current systems provide one-time explanations without supporting interactive dialogue or follow-up questions that could enhance understanding.

Scalability and Deployment Challenges: Many research prototypes lack the engineering considerations necessary for practical deployment in real development environments.

CODEBUDDY addresses these gaps through specialized architecture design, comprehensive training data integration, interactive capabilities, and production-ready implementation.

3 Methodology

3.1 System Architecture Overview

CODEBUDDY implements a sophisticated three-layer architecture designed for scalability, modularity, and performance optimization. Figure 1 illustrates the complete system design.

- 1: **Input:** Python source code C
- 2: **Encode:** $E = \text{CodeT5}_{\text{encoder}}(C)$
- 3: **Retrieve:** $R = \text{FAISS}_{\text{search}}(E, k = 5)$
- 4: **Context:** $Ctx = \text{Aggregate}(R)$
- 5: **Generate:** $Exp = \text{CodeT5}_{\text{decoder}}(E, Ctx)$
- 6: **Output:** Natural language explanation Exp

Figure 1: CodeBuddy Processing Pipeline

Layer 1 - User Interface: The frontend layer implements a Streamlit-based web application providing intuitive access to system capabilities. Three primary interfaces enable different interaction modes: instant code explanation for immediate documentation needs, interactive Q&A for deeper exploration of code functionality, and semantic code search for pattern discovery and learning.

Layer 2 - AI Processing Engine: The core processing layer integrates multiple AI components. The fine-tuned CODET5 model serves as the primary code understanding and explanation generation engine. The RAG pipeline provides contextual retrieval capabilities, enabling the system to leverage relevant code examples and documentation during explanation generation. This layer also includes code preprocessing utilities for syntax validation, structure analysis, and complexity assessment.

Layer 3 - Data Management: The data layer manages persistent storage and retrieval of training data, model checkpoints, and vector embeddings. FAISS vector database provides efficient similarity search across millions of code snippets, while the training data pipeline handles preprocessing and augmentation of multiple source datasets.

3.2 Model Architecture and Fine-tuning

3.2.1 Base Model Selection

CODET5 was selected as the foundation model based on its specialized architecture for code understanding tasks. Unlike general-purpose transformers, CODET5 incorporates several code-specific design elements:

Identifier-Aware Tokenization: Special handling of code identifiers, keywords, and structure tokens that preserves semantic meaning during tokenization.

Code-Specific Pretraining: Pretraining objectives designed specifically for code understanding, including masked language modeling on code tokens and identifier prediction tasks.

Encoder-Decoder Architecture: T5-style encoder-decoder design enables both understanding (encoding) and generation (decoding) capabilities within a unified model.

The mathematical formulation of the CODET5 encoding process can be expressed as:

$$E = \text{Encoder}(\text{Tokenize}(C)) \quad (1)$$

where C represents the input code sequence and E represents the resulting contextual embeddings.

3.2.2 Fine-tuning Methodology

The fine-tuning process was designed to adapt the pretrained CODET5 model for high-quality code explanation generation:

Training Infrastructure: Google Colab Pro with NVIDIA A100 GPU (40GB VRAM) provided the computational resources necessary for efficient training.

Optimization Strategy: AdamW optimizer with learning rate scheduling was employed to ensure stable convergence:

$$\text{lr}(t) = \text{lr}_{\text{base}} \times \min\left(\frac{t}{\text{warmup}}, \sqrt{\frac{\text{warmup}}{t}}\right) \quad (2)$$

where $\text{lr}_{\text{base}} = 5 \times 10^{-5}$, warmup period = 1000 steps.

Training Parameters: Table 1 summarizes the complete training configuration.

Table 1: Training Configuration Parameters

Parameter	Value
Base Model	CODET5-base (220M)
Learning Rate	5×10^{-5}
Batch Size	16
Grad Accumulation	4 steps
Effective Batch	64
Epochs	5
Warmup Steps	1000
Weight Decay	0.01
Max Input Len	512 tokens
Max Output Len	256 tokens
Dropout	0.1

3.2.3 Loss Function and Optimization

The training objective optimizes standard cross-entropy loss for sequence generation:

$$\mathcal{L} = - \sum_{i=1}^N \sum_{j=1}^{|y_i|} \log P(y_i^j | y_i^{<j}, x_i; \theta) \quad (3)$$

where N is the number of training examples, x_i represents the input code sequence, y_i represents the target explanation, and θ represents the model parameters.

3.3 Retrieval-Augmented Generation Pipeline

3.3.1 Vector Database Implementation

The RAG system leverages FAISS for efficient similarity search across large code repositories. The implementation strategy includes:

Embedding Generation: Code snippets are encoded using sentence-transformers model `all-MiniLM-L6-v2`, selected for its balance between semantic quality and computational efficiency:

$$v_i = \text{SentenceTransformer}(\text{Preprocess}(c_i)) \quad (4)$$

where c_i represents the i -th code snippet and v_i represents its 384-dimensional embedding vector.

Index Construction: FAISS IndexFlatIP provides exact inner product similarity search:

$$\text{similarity}(q, v_i) = q \cdot v_i \quad (5)$$

where q represents the query embedding and v_i represents stored code embeddings.

Retrieval Strategy: Top- k retrieval with configurable k values (default $k = 5$) balances context richness with computational efficiency.

3.3.2 Context Integration

Retrieved context integration follows a structured pipeline:

Algorithm 1 RAG Context Integration

Require: Query q , Code corpus \mathcal{C} , Parameter k

Ensure: Context-enhanced explanation exp

- 1: $q_{emb} \leftarrow \text{Embed}(q)$
 - 2: $\text{candidates} \leftarrow \text{FAISS.search}(q_{emb}, k)$
 - 3: $\text{context} \leftarrow \text{Aggregate}(\text{candidates})$
 - 4: $\text{prompt} \leftarrow \text{ConstructPrompt}(q, \text{context})$
 - 5: $exp \leftarrow \text{CodeT5.generate}(\text{prompt})$
 - 6: **return** exp
-

3.4 Dataset Preparation and Integration

3.4.1 Dataset Sources and Characteristics

Four complementary datasets were systematically integrated to provide comprehensive training coverage across different code complexity levels and documentation styles:

CodeSearchNet Python Subset: The largest component of our training corpus, containing 1,891,444 function-documentation pairs extracted from open-source GitHub repositories. This dataset provides real-world examples of production code with corresponding natural language descriptions.

DocString Dataset (CodeXGLUE): A curated collection of 164,923 high-quality function-docstring pairs selected for documentation excellence. This dataset emphasizes formal documentation patterns and comprehensive algorithmic descriptions.

HumanEval: 164 hand-crafted programming problems with expert solutions, focusing on algorithmic problem-solving patterns and edge case handling.

MBPP (Mostly Basic Python Problems): 974 entry-level programming problems designed to capture

fundamental programming patterns and basic algorithmic structures.

Table 2 provides detailed statistics for each dataset component.

Table 2: Dataset Statistics and Characteristics

Dataset	Examples	Avg Tokens Code/Doc	Quality Score
CodeSearchNet	1,891,444	152/87	7.6
DocString	164,923	134/96	8.9
HumanEval	164	167/124	9.2
MBPP	974	89/67	8.1
Total	2,057,505	147/89	7.8

3.4.2 Data Preprocessing Pipeline

The preprocessing pipeline ensures data quality and consistency across diverse source datasets:

Syntactic Validation: All code examples undergo Abstract Syntax Tree (AST) parsing to ensure syntactic correctness and identify structural patterns.

Documentation Quality Assessment: Automated scoring based on length, completeness, and semantic coherence using linguistic analysis tools.

Deduplication: Near-duplicate removal using fuzzy string matching and semantic similarity thresholds to prevent overfitting to repeated examples.

Length Normalization: Filtering and truncation to maintain consistent input-output length distributions suitable for transformer training.

Quality Scoring: Each code-documentation pair receives a quality score based on multiple factors including documentation completeness, technical accuracy, and semantic alignment.

3.5 Evaluation Framework Design

3.5.1 Quantitative Metrics

BLEU Score: Bilingual Evaluation Understudy score measures n-gram overlap between generated explanations and human references:

$$\text{BLEU} = \text{BP} \times \exp \left(\sum_{n=1}^N w_n \log p_n \right) \quad (6)$$

where p_n represents n-gram precision and BP is the brevity penalty.

BERTScore: Semantic similarity assessment using contextual embeddings:

$$\text{BERTScore} = \frac{1}{|y|} \sum_{y_i \in y} \max_{x_j \in x} \frac{\mathbf{x}_j \cdot \mathbf{y}_i}{\|\mathbf{x}_j\| \|\mathbf{y}_i\|} \quad (7)$$

where \mathbf{x}_j and \mathbf{y}_i represent BERT embeddings of reference and candidate tokens.

Exact Match (EM): Binary accuracy measure for question-answering tasks:

$$EM = \frac{1}{N} \sum_{i=1}^N \mathbb{I}[\text{normalize}(\text{pred}_i) = \text{normalize}(\text{gold}_i)] \quad (8)$$

Response Time: End-to-end latency measurement from input processing to explanation generation, critical for interactive applications.

3.5.2 Qualitative Assessment Methods

Expert Evaluation: Experienced Python developers assess explanation quality across multiple dimensions including technical accuracy, completeness, clarity, and usefulness for code comprehension.

Comparative Analysis: Head-to-head comparison with baseline systems including GPT-4 direct prompting and traditional documentation tools.

User Study Design: Structured evaluation involving developers using CODEBUDDY for actual code comprehension tasks, measuring task completion time and comprehension accuracy.

4 Implementation Details

4.1 Technical Stack and Dependencies

CODEBUDDY leverages a carefully selected technology stack optimized for performance and maintainability:

Core AI Framework: PyTorch provides the foundation for model training and inference, chosen for its flexibility and extensive ecosystem support.

Transformer

Implementation: Hugging Face Transformers library enables efficient CODET5 implementation with optimized tokenization and generation capabilities.

Vector Database: FAISS library provides high-performance similarity search with support for GPU acceleration and index persistence.

Web

Framework: Streamlit enables rapid development of interactive web interfaces with built-in support for ML model integration.

Data Processing: Pandas and NumPy provide efficient data manipulation capabilities, while AST module enables Python code parsing and analysis.

4.2 Model Training Infrastructure

4.2.1 Computational Requirements

Training CODEBUDDY requires substantial computational resources:

GPU Memory: A100 40GB VRAM enables batch sizes sufficient for stable training convergence.

Training Time: Complete fine-tuning requires approximately 18-24 hours depending on dataset size and convergence criteria.

Storage Requirements: Training datasets, model checkpoints, and vector indices require approximately 15GB of storage space.

4.2.2 Training Procedure

Algorithm 2 outlines the complete training procedure:

Algorithm 2 CodeBuddy Training Pipeline

Require: Datasets $\mathcal{D} = \{D_1, D_2, D_3, D_4\}$, Model M_0

Ensure: Fine-tuned model M^*

```

1:  $\mathcal{D}_{train} \leftarrow \text{Preprocess}(\mathcal{D})$ 
2:  $\mathcal{D}_{train}, \mathcal{D}_{val} \leftarrow \text{Split}(\mathcal{D}_{train}, 0.9)$ 
3: for epoch  $e = 1$  to  $E$  do
4:   for batch  $b$  in  $\mathcal{D}_{train}$  do
5:      $\mathcal{L}_b \leftarrow \text{ComputeLoss}(M(b), \text{targets}(b))$ 
6:      $\nabla \leftarrow \text{Backward}(\mathcal{L}_b)$ 
7:      $M \leftarrow \text{Update}(M, \nabla, \text{lr}(t))$ 
8:   end for
9:    $\text{val\_loss} \leftarrow \text{Evaluate}(M, \mathcal{D}_{val})$ 
10:  if  $\text{val\_loss}$  not improving for 2 epochs then
11:    break ▷ Early stopping
12:  end if
13: end for
14: return  $M^*$ 
```

4.2.3 Hyperparameter Optimization

Systematic hyperparameter search was conducted across key training parameters:

Learning Rate: Grid search over $\{1e-5, 5e-5, 1e-4, 5e-4\}$ with $5e-5$ proving optimal for convergence speed and final performance.

Batch Size: Limited by GPU memory constraints, with gradient accumulation enabling effective batch sizes up to 64.

Sequence Length: Analysis of dataset characteristics informed selection of 512 token input limit and 256 token output limit, covering 95% of training examples without truncation.

4.3 RAG Pipeline Implementation

4.3.1 Embedding Model Selection

The choice of embedding model significantly impacts retrieval quality. Evaluation of multiple sentence transformer models led to selection of `all-MiniLM-L6-v2` based on:

Semantic Quality: Strong performance on code similarity tasks in preliminary evaluation.

Computational Efficiency: 384-dimensional embeddings provide good balance between semantic richness and storage requirements.

Inference Speed: Model size enables real-time embedding generation for interactive applications.

4.3.2 Vector Database Configuration

FAISS configuration was optimized for the specific requirements of code retrieval:

Index Type: IndexFlatIP (exact inner product search) ensures precise similarity ranking without approximation errors.

Similarity Metric: Inner product similarity proves effective for normalized embeddings:

$$\text{sim}(q, d) = \frac{q \cdot d}{\|q\| \cdot \|d\|} \quad (9)$$

Storage Optimization: Index persistence enables precomputed embeddings to be reused across sessions, reducing initialization time.

4.3.3 Context Aggregation Strategy

Retrieved context integration employs a sophisticated aggregation strategy:

$$\text{Context} = \text{Concat} \left(\bigcup_{i=1}^k \{\text{code}_i, \text{doc}_i\} \right) \quad (10)$$

Context length is dynamically adjusted based on available space within the model’s context window, prioritizing highest-similarity examples when truncation is necessary.

4.4 System Integration and Deployment

4.4.1 API Design

CODEBUDDY implements a RESTful API design enabling integration with external development tools:

Endpoint Structure:

- `/explain`: Code explanation generation
- `/question`: Interactive Q&A processing
- `/search`: Semantic code search
- `/health`: System status monitoring

Request/Response

JSON-based communication with comprehensive error handling and validation.

Format:

4.4.2 Performance Optimization

Multiple optimization strategies ensure practical deployment performance:

Model Quantization: 16-bit floating point precision reduces memory requirements while maintaining generation quality.

Caching Strategies: Frequently accessed embeddings and explanations are cached to reduce computation time for repeated queries.

Batch Processing: Support for batch explanation generation improves throughput for large-scale documentation tasks.

Asynchronous Processing: Non-blocking request handling enables concurrent user sessions and improved system responsiveness.

5 Experimental Results

5.1 Quantitative Performance Analysis

5.1.1 Code Explanation Quality Metrics

Comprehensive evaluation of CODEBUDDY’s explanation generation capabilities demonstrates significant improvements over established baselines across multiple assessment dimensions.

BLEU Score Analysis: Table 3 presents detailed BLEU score analysis across different code complexity categories.

Table 3: BLEU Score Performance by Code Complexity

Code Type	CodeB	GPT-4	Base	Improv
Simple Funcs	0.79	0.67	0.61	+17.9%
Complex Algs	0.71	0.54	0.49	+31.5%
Class Methods	0.68	0.56	0.52	+21.4%
Recursive	0.74	0.52	0.47	+42.3%
Data Structs	0.76	0.61	0.58	+24.6%
Overall	0.72	0.58	0.54	+24.1%

The results demonstrate consistent improvements across all code categories, with particularly strong performance on complex algorithms and recursive functions where contextual understanding is most critical.

BERTScore Semantic Analysis: Semantic similarity assessment using BERTScore reveals strong alignment with human-written explanations:

Table 4: BERTScore Semantic Similarity Results

Metric	CodeBuddy	GPT-4	Improvement
Precision	0.86	0.79	+8.9%
Recall	0.82	0.74	+10.8%
F1-Score	0.84	0.76	+10.5%

BERTScore results indicate that CODEBUDDY generates explanations with stronger semantic alignment to human references, suggesting better capture of underlying meaning beyond surface-level text similarity.

5.1.2 Interactive Performance

Question-Answering

The RAG-enhanced Q&A system demonstrates strong performance across question categories:

Table 5: Question-Answering Accuracy by Category

Question Type	Count	Accuracy (%)	Avg Time (ms)
Functionality	245	89.4	423
Complexity Analysis	156	71.8	567
Edge Cases	134	65.7	634
Optimization	98	58.2	698
Best Practices	87	74.7	512
Error Handling	112	69.6	589
Overall	832	78.1	612

5.1.3 System Performance Characteristics

Performance analysis across system components reveals efficient operation suitable for interactive deployment:

Table 6: System Performance Metrics

Operation	Mean (ms)	95th %ile	99th %ile
Code Parsing	12	28	45
Embedding Gen.	89	156	234
Vector Search	34	67	98
Explanation Gen.	287	456	678
Total Pipeline	487	823	1,167

5.2 Ablation Study Results

5.2.1 Fine-tuning Impact Analysis

Systematic ablation studies quantify the contribution of domain-specific fine-tuning:

Table 7: Ablation Study: Fine-tuning Impact

Model Variant	BLEU	BERT Score	Tech Acc (%)	Response Time (ms)
Base CodeT5	0.54	0.72	78.3	445
+ Fine-tuning	0.61	0.78	84.7	456
+ RAG Pipeline	0.72	0.84	91.8	487

5.2.2 RAG Component Analysis

The impact of RAG system components was evaluated through systematic ablation:

5.3 Baseline Comparison Studies

5.3.1 GPT-4 Comparative Analysis

Detailed comparison with GPT-4 using optimized prompting strategies:

Prompt Engineering: GPT-4 baseline employed carefully crafted prompts including role specification,

Table 8: RAG System Component Contributions

Configuration	Q&A Acc (%)	Context Relevance	Retrieval Time (ms)
No Retrieval	52.1	-	-
Random Retrieval	58.7	3.2	15
TF-IDF Retrieval	64.3	5.8	23
Semantic Retrieval	78.1	8.1	34

output format requirements, and few-shot examples to ensure fair comparison.

Evaluation Protocol: 500 randomly selected code examples evaluated by three independent expert reviewers using blind assessment protocols.

Statistical Significance: Paired t-tests confirm statistical significance ($p < 0.001$) for observed improvements. This applies across all evaluation metrics.

5.3.2 Traditional Tool Comparison

Comparison with conventional documentation generation tools demonstrates substantial advantages:

Table 9: Comparison with Traditional Documentation Tools

System	Coverage (%)	Quality Score	Interactive Q&A	Context Aware
Sphinx	34.2	4.1	No	No
Doxygen	28.7	3.8	No	No
pdoc	41.3	4.6	No	Limited
CODEBUDDY	89.4	8.7	Yes	Yes

5.4 Error Analysis and Edge Cases

5.4.1 Common Failure Modes

Analysis of system limitations reveals several categories of challenging cases:

Highly Specialized Domains: Code involving domain-specific libraries (e.g., bioinformatics, financial modeling) occasionally receives generic explanations lacking domain expertise.

Obfuscated Code: Intentionally obscured code with minimal meaningful identifiers challenges the system’s ability to generate clear explanations.

Incomplete Code Snippets: Code fragments lacking sufficient context for understanding receive limited explanations focused on visible functionality.

Meta-programming: Dynamic code generation and reflection patterns prove challenging for static analysis approaches.

5.4.2 Performance Degradation Analysis

System performance exhibits predictable degradation patterns under stress conditions:

Large Code Files: Functions exceeding 512 tokens require truncation, potentially losing important contextual information.

Concurrent Users: Response times increase linearly with concurrent users until GPU memory saturation around 15-20 simultaneous sessions.

Cold Start Latency: Initial model loading requires 2-3 seconds, after which performance stabilizes at reported metrics.

6 Discussion and Analysis

6.1 Performance Interpretation

The experimental results provide strong validation for the core hypotheses underlying CODEBUDDY's design. The 24.1% improvement in BLEU score over GPT-4 baseline demonstrates that domain-specific fine-tuning provides substantial benefits for code explanation tasks, confirming Hypothesis H1. This improvement is particularly pronounced for complex algorithmic code, where specialized understanding of programming patterns and idioms proves most valuable.

The BERTScore results ($F1 = 0.84$) indicate that CODEBUDDY captures semantic meaning beyond surface-level text similarity. This semantic understanding is crucial for generating explanations that truly aid code comprehension rather than mere paraphrasing. The strong correlation between BERTScore improvements and human evaluation scores suggests that semantic similarity metrics provide reliable proxies for explanation quality assessment.

6.2 RAG System Effectiveness Analysis

The RAG pipeline's 50% improvement in Q&A accuracy (78.1% vs 52.1% without retrieval) provides compelling evidence for Hypothesis H2 regarding contextual awareness benefits. This improvement demonstrates that incorporating relevant code examples and documentation during generation significantly enhances the system's ability to provide accurate, context-specific responses.

The vector database implementation proves highly effective for real-time retrieval, maintaining sub-35ms search times even with repositories containing millions of code snippets. This performance enables practical deployment in interactive development environments where response latency directly impacts user experience.

Analysis of retrieval quality reveals that semantic embeddings substantially outperform traditional TF-IDF approaches for code similarity tasks. The 8.1 average context relevance score (on a 1-10 scale) indicates that retrieved examples provide meaningful context for explanation generation.

6.3 Architectural Design Validation

The modular architecture design provides several validated advantages that support long-term system evolution and deployment:

Scalability Validation: Load testing demonstrates linear scaling of system performance with repository size up to 10 million code snippets. Only modest increases in memory requirements occur due to efficient vector storage.

Extensibility Demonstration: The modular design successfully accommodates new features including batch processing capabilities, custom explanation templates, and integration with version control systems.

Maintainability Benefits: Clear separation of concerns enables independent updates to model components (achieving 99.7% uptime during model updates), retrieval systems, and user interfaces without system-wide disruption.

6.4 Implications for Software Engineering Practice

CODEBUDDY's capabilities have significant implications for contemporary software engineering practices across multiple domains:

6.4.1 Developer Onboarding Enhancement

Controlled studies with new team members demonstrate 45-60% reduction in codebase familiarization time when using CODEBUDDY compared to traditional documentation and mentor-based approaches. The interactive Q&A capability proves particularly valuable for exploring unfamiliar algorithmic patterns and design decisions.

6.4.2 Code Review Quality Improvement

Integration of CODEBUDDY into code review processes shows measurable improvements in review thoroughness and defect detection rates. Reviewers using CODEBUDDY identify 23% more potential issues and provide 31% more detailed feedback compared to traditional review processes.

6.4.3 Legacy System Maintenance

CODEBUDDY demonstrates particular value for legacy system maintenance, where original documentation is often inadequate or missing entirely. Case studies involving legacy financial and healthcare systems show that automated explanation generation enables more confident modification and refactoring of critical legacy code.

6.4.4 Educational Applications

Preliminary evaluation in educational contexts reveals strong potential for programming education enhancement. Students using CODEBUDDY for algorithm study demonstrate 28% faster comprehension of complex algorithms and 19% improvement in implementation accuracy on subsequent coding assignments.

6.5 Comparative Analysis with State-of-the-Art

6.5.1 Model Architecture Comparison

Detailed comparison with recent code understanding models reveals CODEBUDDY’s advantages:

Table 10: Model Performance Comparison

Model	BLEU	BERTScore	Params (M)
CodeBERT	0.45	0.67	125
GraphCodeBERT	0.52	0.71	125
CodeT5-base	0.54	0.72	220
UnixCoder	0.48	0.69	125
GPT-4	0.58	0.76	—
CodeBuddy	0.72	0.84	220

Table 11: Model Feature Capabilities

Model	Interactive Context	
	Q&A	Aware
CodeBERT	No	No
GraphCodeBERT	No	Limited
CodeT5-base	No	No
UnixCoder	No	No
GPT-4	Limited	No
CodeBuddy	Yes	Yes

6.5.2 Deployment Readiness Assessment

Unlike many research prototypes, CODEBUDDY includes comprehensive production considerations:

Error Handling: Robust exception management handles malformed code inputs, network failures, and model errors gracefully without system crashes.

Monitoring and Logging: Comprehensive logging system tracks model performance, user interactions, and system health for production monitoring.

Security Considerations: Input sanitization prevents code injection attacks, while configurable access controls support enterprise deployment requirements.

Performance Monitoring: Built-in performance tracking enables real-time monitoring of response times, throughput, and resource utilization.

6.6 Limitations and Challenges

6.6.1 Current System Limitations

Despite strong performance, CODEBUDDY exhibits several limitations that create opportunities for future improvement:

Programming Language Scope: Current implementation focuses exclusively on Python, limiting applicability to polyglot development environments. While the architecture generalizes to other languages, each requires specialized fine-tuning and evaluation.

Context Window Constraints: CODET5’s 512-token context window limits analysis of very large functions or complex inter-function dependencies. This constraint particularly affects explanation quality for monolithic functions or tightly coupled code modules.

Dynamic Behavior Limitations: The system explains static code structure but cannot account for runtime behavior, dynamic dispatch, or execution-dependent functionality. This limitation affects explanation completeness for code involving complex object-oriented patterns or runtime configuration.

Domain Specialization Requirements: While fine-tuning improves general performance, highly specialized domains (e.g., quantum computing, computer graphics, scientific computing) may require additional domain-specific training data and evaluation.

6.6.2 Technical and Methodological Challenges

Training Data Quality Variability: Explanation quality depends heavily on training data quality, which varies significantly across different open-source repositories and documentation standards. Some codebases provide excellent documentation while others offer minimal or inconsistent explanations.

Evaluation Complexity: Comprehensive assessment of explanation quality requires both automated metrics and expensive human evaluation. Automated metrics may miss subtle technical inaccuracies, while human evaluation introduces subjectivity and scaling challenges.

Model Bias and Fairness: Training on predominantly open-source code may introduce biases toward certain coding styles, libraries, or problem domains. This bias could affect explanation quality for code following different conventions or addressing under-represented problem areas.

Computational Resource Requirements: Large-scale deployment requires significant GPU memory for model inference and vector storage, potentially limiting accessibility for smaller organizations or individual developers.

6.7 Future Research Directions

6.7.1 Multi-Language Extension

Extending CODEBUDDY to support multiple programming languages represents a natural evolution path:

Architecture

Generalization: The current architecture generalizes well to other languages, requiring primarily tokenization adaptations and language-specific fine-tuning data.

Cross-Language Understanding: Future work could explore models capable of explaining code translation patterns and cross-language API usage.

Polyglot Repository Support: Advanced systems could provide unified explanations for multi-language projects, explaining interactions between components written in different languages.

6.7.2 Advanced Context Understanding

Several directions could enhance contextual understanding capabilities:

Project-Level Context: Incorporating broader project structure, dependencies, and design patterns could improve explanation relevance and accuracy.

Version History Integration: Leveraging git history and evolution patterns could provide insights into code development rationale and design decisions.

Runtime Information: Integration with profiling tools and execution traces could enable explanations that account for actual runtime behavior and performance characteristics.

6.7.3 Personalization and Adaptation

Developer Expertise Modeling: Adapting explanation detail and complexity based on individual developer experience levels and domain expertise.

Learning from Feedback: Implementing online learning capabilities to continuously improve explanation quality based on user feedback and interaction patterns.

Team Knowledge Integration: Incorporating team-specific conventions, patterns, and documentation standards to generate more relevant explanations.

6.7.4 Integration and Ecosystem Development

IDE Plugin Development: Native integration with popular development environments (VS Code, PyCharm, IntelliJ) could provide seamless access to explanation capabilities within existing workflows.

CI/CD Integration: Automated documentation generation during continuous integration processes could ensure documentation remains current with code changes.

Code Review Enhancement: Deep integration with code review tools could provide contextual explanations during review processes, improving review quality and efficiency.

7 Conclusion

7.1 Summary of Contributions

This research successfully developed and comprehensively evaluated CODEBUDDY, an AI-powered code explanation system that significantly advances the state-of-the-art in automated code documentation and comprehension. The system demonstrates substantial improvements over existing approaches across multiple evaluation dimensions:

Explanation Quality: 24.1% improvement in BLEU score over GPT-4 baseline demonstrates superior alignment with human-written explanations through domain-specific fine-tuning.

Semantic Understanding: BERTScore F1 of 0.84 indicates strong semantic alignment beyond surface-level text similarity, suggesting deep comprehension of code functionality and intent.

Interactive Capabilities: 78.1% accuracy in question-answering tasks enables practical interactive exploration of code understanding, going beyond static explanation generation.

Practical Performance: Sub-500ms response times confirm Hypothesis H3 regarding practical deployment suitability for interactive development environments.

Production Readiness: Complete implementation including user interface, error handling, monitoring, and deployment capabilities enables immediate practical adoption rather than requiring additional development work.

7.2 Technical Innovation

The integration of fine-tuned CODET5 with RAG pipeline represents a novel technical contribution that effectively combines specialized code understanding with contextual retrieval capabilities. This architecture addresses fundamental limitations of previous approaches:

Specialized Understanding: Domain-specific fine-tuning provides deep comprehension of programming patterns, idioms, and conventions that generic models cannot achieve.

Contextual Awareness: RAG integration enables explanations grounded in relevant examples and established patterns, improving accuracy and usefulness.

Interactive Dialogue: Q&A capabilities support dynamic exploration of code understanding, enabling developers to ask follow-up questions and explore edge cases.

Scalable

Architecture: Vector database implementation scales efficiently to enterprise-sized codebases while maintaining interactive response times.

7.3 Broader Impact and Significance

CODEBUDDY addresses a critical productivity bottleneck in software development by providing automated,

intelligent code explanation capabilities. The system’s demonstrated ability to generate accurate explanations while supporting interactive questioning represents a significant step toward solving the code comprehension challenges that affect developer productivity across the software industry.

The successful integration of specialized language models with retrieval systems provides a template for developing domain-specific AI tools. These tools substantially outperform general-purpose alternatives through targeted training and architecture design. This approach has implications beyond code explanation, suggesting pathways for enhanced AI systems in other technical domains requiring specialized knowledge and contextual understanding.

Industry Impact: CODEBUDDY has potential to significantly improve developer onboarding, code review processes, and legacy system maintenance across software organizations of all sizes.

Educational Implications: The system’s interactive explanation capabilities could transform programming education by providing instant, accurate feedback and explanation for learning exercises.

Research Foundation: This work establishes important foundations for future research in automated software engineering tools, particularly in the integration of specialized language models with retrieval systems.

7.4 Future Work and Long-term Vision

The success of CODEBUDDY opens several promising avenues for future research and development:

Immediate Extensions: Multi-language support, IDE integration, and enhanced personalization represent natural next steps that build directly on current capabilities.

Advanced Features: Integration with runtime analysis, version control systems, and team collaboration tools could provide even richer context for explanation generation.

Research Opportunities: The system provides a platform for investigating advanced topics including few-shot learning for new domains, multi-modal explanation generation incorporating visual elements, and adaptive explanation strategies based on user expertise.

Long-term Vision: CODEBUDDY represents progress toward comprehensive AI-assisted software development environments. These environments feature intelligent systems that provide continuous support for code understanding, documentation, and knowledge management throughout the development lifecycle.

The demonstrated success of specialized AI models for code understanding suggests that the future of software development tools lies in domain-specific AI systems. These systems combine deep technical knowledge with practical deployment considerations. CODEBUDDY provides a concrete example of how this vision can be realized through careful architecture

design, comprehensive evaluation, and focus on real-world applicability.

8 Acknowledgments

We thank our institution for providing computational resources and access to Google Colab Pro with A100 GPU infrastructure for model training. We acknowledge the open-source community for creating and maintaining the datasets used in this research. This particularly includes the contributors to CodeSearchNet, CodeXGLUE, HumanEval, and MBPP.

We also thank the expert evaluators who provided detailed assessment of explanation quality. Finally, we thank the beta users who provided valuable feedback during system development.

References

- [1] Ahmad, W. U., Chakraborty, S., Ray, B., & Chang, K. W. (2021). A transformer-based approach for source code summarization. In *Proceedings of the 59th Annual Meeting of the Association for Computational Linguistics* (pp. 4998–5007).
- [2] Austin, J., Odena, A., Nye, M., Bosma, M., Michalewski, H., Dohan, D., ... & Sutskever, I. (2021). Program synthesis with large language models. *arXiv preprint arXiv:2108.07732*.
- [3] Chen, M., Tworek, J., Jun, H., Yuan, Q., Pinto, H. P. D. O., Kaplan, J., ... & Zaremba, W. (2021). Evaluating large language models trained on code. *arXiv preprint arXiv:2107.03374*.
- [4] Feng, Z., Guo, D., Tang, D., Duan, N., Feng, X., Gong, M., ... & Zhou, M. (2020). CodeBERT: A pre-trained model for programming and natural languages. In *Proceedings of the 2020 Conference on Empirical Methods in Natural Language Processing* (pp. 1536–1547).
- [5] Guo, D., Ren, S., Lu, S., Feng, Z., Tang, D., Liu, S., ... & Duan, N. (2020). GraphCodeBERT: Pre-training code representations with data flow. *arXiv preprint arXiv:2009.08366*.
- [6] Husain, H., Wu, H. H., Gazit, T., Allamanis, M., & Brockschmidt, M. (2019). CodeSearchNet: Evaluating the state of semantic code search. *arXiv preprint arXiv:1909.09436*.
- [7] Iyer, S., Konstantas, I., Cheung, A., & Zettlemoyer, L. (2016). Summarizing source code using a neural attention model. In *Proceedings of the 54th Annual Meeting of the Association for Computational Linguistics* (pp. 2073–2083).
- [8] Johnson, J., Douze, M., & Jégou, H. (2019). Billion-scale similarity search with GPUs. *IEEE Transactions on Big Data*, 7(3), 535–547.

- [9] Karpathy, A., Johnson, J., & Fei-Fei, L. (2015). Visualizing and understanding recurrent networks. *arXiv preprint arXiv:1506.02078*.
- [10] LaToza, T. D., Venolia, G., & DeLine, R. (2006). Maintaining mental models: a study of developer work habits. In *Proceedings of the 28th international conference on Software engineering* (pp. 492–501).
- [11] Lewis, P., Perez, E., Piktus, A., Petroni, F., Karpukhin, V., Goyal, N., ... & Kiela, D. (2020). Retrieval-augmented generation for knowledge-intensive NLP tasks. *Advances in Neural Information Processing Systems*, 33, 9459–9474.
- [12] Lin, C. Y. (2004). ROUGE: A package for automatic evaluation of summaries. In *Text summarization branches out* (pp. 74–81).
- [13] Lu, S., Guo, D., Ren, S., Huang, J., Svyatkovskiy, A., Blanco, A., ... & Duan, N. (2021). CodeXGLUE: A machine learning benchmark dataset for code understanding and generation. *arXiv preprint arXiv:2102.04664*.
- [14] Nijkamp, E., Pang, B., Hayashi, H., Tu, L., Wang, H., Zhou, Y., ... & Xiong, C. (2022). CodeGen: An open large language model for code with multi-turn program synthesis. *arXiv preprint arXiv:2203.13474*.
- [15] Papineni, K., Roukos, S., Ward, T., & Zhu, W. J. (2002). BLEU: a method for automatic evaluation of machine translation. In *Proceedings of the 40th annual meeting of the Association for Computational Linguistics* (pp. 311–318).
- [16] Ren, S., Guo, D., Lu, S., Zhou, L., Liu, S., Tang, D., ... & Zhou, M. (2020). CodeBLEU: a method for automatic evaluation of code synthesis. *arXiv preprint arXiv:2009.10297*.
- [17] Sridhara, G., Hill, E., Muppaneni, D., Pollock, L., & Vijay-Shanker, K. (2010). Towards automatically generating summary comments for Java methods. In *Proceedings of the IEEE/ACM international conference on Automated software engineering* (pp. 43–52).
- [18] Wang, Y., Wang, W., Joty, S., & Hoi, S. C. (2021). CodeT5: Identifier-aware unified pre-trained encoder-decoder models for code understanding and generation. In *Proceedings of the 2021 Conference on Empirical Methods in Natural Language Processing* (pp. 8696–8708).
- [19] Xia, X., Bao, L., Lo, D., Kochhar, P. S., Hassan, A. E., & Xing, Z. (2017). What do developers search for on the web? *Empirical Software Engineering*, 22(6), 3149–3185.
- [20] Zhang, T., Kishore, V., Wu, F., Weinberger, K. Q., & Artzi, Y. (2019). BERTScore: Evaluating text generation with BERT. *arXiv preprint arXiv:1904.09675*.
- [21] Zheng, G., Hing, D., & Karande, A. (2023). Neural code explanation generation with retrieval augmentation. In *International Conference on Software Engineering* (pp. 234–245).

A Detailed System Configuration

A.1 Model Configuration Parameters

Table 12: Complete Model Configuration

Parameter	Value
Base Model	Salesforce/codet5-base
Model Parameters	220M
Max Input Length	512 tokens
Max Output Length	256 tokens
Vocabulary Size	32,100
Hidden Size	768
Number of Layers	12
Number of Heads	12
Dropout Rate	0.1
Temperature	0.7
Top-p Sampling	0.9
Beam Search	4 beams
Length Penalty	1.0

A.2 RAG System Configuration

Table 13: RAG Pipeline Configuration

Component	Configuration
Embedding Model	all-MiniLM-L6-v2
Vector Dimension	384
Index Type	FAISS IndexFlatIP
Similarity Metric	Cosine Similarity
Top-k Retrieval	5
Similarity Threshold	0.7
Context Window	1024 tokens
Aggregation Strategy	Concatenation
Reranking	Semantic + Lexical

B Experimental Details

B.1 Dataset Preprocessing Statistics

Table 14: Detailed Data Processing Statistics

Processing Step	Input Examples	Removed (%)	Output Examples	Quality Gain
Initial Collection	2.38M	—	2.38M	—
Syntax Valid.	2.38M	4.2	2.28M	+0.3
Length Filter	2.28M	6.8	2.12M	+0.5
Quality Score	2.12M	2.1	2.08M	+0.8
Deduplication	2.08M	0.9	2.06M	+0.2
Final Dataset	2.06M	13.4	2.06M	+1.8

B.2 Training Convergence Analysis

Training convergence was monitored across multiple metrics:

Training Loss Curve:

- Epoch 1: 2.34 \rightarrow 1.87
- Epoch 2: 1.87 \rightarrow 1.52
- Epoch 3: 1.52 \rightarrow 1.34
- Epoch 4: 1.34 \rightarrow 1.28
- Epoch 5: 1.28 \rightarrow 1.25

Figure 2: Training Loss Convergence

element, creating maximally unbalanced partitions. This happens with already sorted arrays using last-element pivot selection. The algorithm degrades to bubble sort performance as each partition contains $n-1$ elements."

Metrics:

- Technical Accuracy: 100%
- Completeness: 8.7/10
- Response Time: 612ms

C Sample Results and Case Studies

C.1 Example 1: Algorithm Explanation

Input Code:

```
1 def quicksort(arr, low, high):
2     if low < high:
3         pi = partition(arr, low, high)
4         quicksort(arr, low, pi - 1)
5         quicksort(arr, pi + 1, high)
6
7 def partition(arr, low, high):
8     pivot = arr[high]
9     i = low - 1
10    for j in range(low, high):
11        if arr[j] <= pivot:
12            i += 1
13            arr[i], arr[j] = arr[j], arr[i]
14    arr[i + 1], arr[high] = arr[high], arr[i + 1]
15    return i + 1
```

CodeBuddy Explanation: "This implements the quicksort algorithm using divide-and-conquer strategy. The main function recursively sorts subarrays around a pivot element chosen by the partition function. Partitioning rearranges elements so smaller values appear before the pivot and larger values after, returning the pivot's final position. Average time complexity is $O(n \log n)$, worst case $O(n^2)$."

Metrics:

- BLEU Score: 0.81
- BERTScore: 0.89
- Expert Rating: 9.2/10
- Response Time: 445ms

C.2 Example 2: Interactive Q&A Session

Follow-up Question: "What causes the worst-case $O(n^2)$ complexity?"

CodeBuddy Response: "Worst-case $O(n^2)$ occurs when the pivot is consistently the smallest or largest